


ASSIGNMENT 0: DSP-Experiments

FINAL PROJECT

AVRUTIN, PERR-SAUER



CS102

Nicolas Avrutin,

2014 EE

nicolasavru@gmail.com

CS102

Jordan Perr-Sauer

2014 EE

jordan@jperr.com

Wednesday 11th August, 2010

Wednesday 11th August, 2010

Jordan Perr-Sauer

Class of 2014

Electrical Engineering

`jordan@jperr.com`

Nicolas Avrutin

Class of 2014

Electrical Engineering

`nicolasavru@gmail.com`

Stack Overflow

`http://stackoverflow.com/questions/2063284/what-is-the-easiest-way-to-read-wav-files-us`

Scott Wilson, Stanford University

`https://ccrma.stanford.edu/courses/422/projects/WaveFormat/`

1 Description

DSP-Experiments consists of python scripts that can encode images into sound files such that the image can be seen in the spectrogram of the produced sound. The code also provides a method for viewing the spectrogram of the generated audio, decoding the sound back into a (lossy) version of the original image.

In our spectrograms, the x-axis represents time (in the audio file) and the y-axis represents frequency. The brightness of a given pixel represents amplitude (loudness). To encode an image such that it can be viewed in a spectrogram, we first separate the image by it's columns of pixels. Each column of pixels will be represented by a “tone” in the output audio stream

that lasts for a fixed amount of time. These tones consist of sin waves of varying frequencies, at varying amplitudes, depending on the intensity and position of each pixel in the column, respectively. Each column in the input image is converted into such a tone, and these tones are concatenated and output to a wave file.

2 Compilation

You do not need to compile any part of this program. There are some dependencies, however, which you must install prior to execution.

2.1 Dependencies

- Python 2.6 or higher (NumPy and PIL are not yet compatible with Python 3)
- NumPy (<http://numpy.scipy.org/>)
- Python Imaging Library (<http://www.pythonware.com/products/pil/>)

3 Execution

3.1 Manual Page

To produce sound (as a wav file) from an image:

```
python imageToWav.py g|c pathToImage pathToWav
```

The g and c options stand for “Grayscale” or “Color”. Using the c option will result in a 3 channel wav file, where each channel of audio represents one channel of color (red, green, and blue). Using the g option will result in a single channel (mono) wav file and all color data will be lost.

To generate the spectrogram of a produced (or other) wav file:

```
python wavToImage.py pathToWav pathToImage
```

The format of the image will be derived from the extension you give pathToImage. You can use any format supported by PIL.

3.2 Sample Inputs

Sample images are provided in the test_images dir:

```
python imageToWav.py c test_images/test.png out.wav
```

Sample wav files generate with our script are located in outputs:

```
python wavToImage.py outputs/cpu_color.wav out.png
```

4 Features

We support color images by using 3-channel wav files. You can specify grayscale or color using command line arguments at execution. This is discussed in “Execution.”

5 Notes

- Scaling and resizing algorithms will be significantly reworked in the future.

6 Listings

6.1 imageToWav.py

```
1  """
2
3  imageToWav.py
4  Encodes image data as an audio signal.
5  One channel of audio per channel of image data.
6
7  Author: Nicolas Avrutin, nicolasavru@gmail.com
8          Jordan Perr-Sauer, jordan@jperr.com
9
10 Revision history:
11     See http://github.com/nicolasavru/DSP-Experiments
12
13 Bugs:
14
15 Todo:
16     - Tweak scaling/encoding algorithm to lessen artifacts
17     - Variable "slice" width for more space efficient encoding
18
19 """
20
21 ##### IMPORTS #####
22
23 import numpy as np
24 import Image, struct, math, sys
25
26
27 ##### DEFS AND ARGS #####
28
29 SAMPLE_RATE = 44100
30
31 YRES = 400
32 T_PER_COL = 0.03
33
34 ARG_IMAGE = sys.argv[2]
35 ARG_OUTFILE = sys.argv[3]
36 ARG_COLOR = False
37 if sys.argv[1] == "c":
38     ARG_COLOR = True
39
40
41 CHANNELS = 1
42 if ARG_COLOR:
43     CHANNELS = 3
44
45 #rgb aliases
46 R=0
47 G=1
48 B=2
49
```

```

50
51 ##### FUNCTIONS #####
52
53 def oscillator(x, freq=1, amp=1, base=0, phase=0):
54     return base + amp * np.sin(2 * np.pi * freq * x + phase)
55
56 def writewav(filename, numChannels, sampleRate, bitsPerSample, nSamples, data):
57     wave = open(filename, 'wb')
58     dataSize = nSamples * numChannels * bitsPerSample / 8
59     #https://ccrma.stanford.edu/courses/422/projects/WaveFormat/
60     ChunkID = 'RIFF'
61     ChunkSize = struct.pack('<I', dataSize + 36)
62     Format = 'WAVE'
63     Subchunk1ID = 'fmt '
64     Subchunk1Size = struct.pack('<I', 16)
65     AudioFormat = struct.pack('<H', 1)
66     NumChannels = struct.pack('<H', numChannels)
67     SampleRate = struct.pack('<I', sampleRate)
68     ByteRate = struct.pack('<I', sampleRate * numChannels * bitsPerSample / 8)
69     BlockAlign = struct.pack('<H', numChannels * bitsPerSample / 8)
70     BitsPerSample = struct.pack('<H', bitsPerSample)
71     Subchunk2ID = 'data'
72     Subchunk2Size = struct.pack('<I', dataSize)
73     header = ChunkID + ChunkSize + Format + Subchunk1ID + Subchunk1Size + \
74             AudioFormat + NumChannels + SampleRate + ByteRate + BlockAlign + \
75             BitsPerSample + Subchunk2ID + Subchunk2Size
76     wave.write(header)
77     # higher amplitude causes noise (vertical bars)
78     print "Packing WAV..."
79     (1000 * data).astype(np.int16).tofile(wave)
80     wave.close()
81
82
83 ##### MAIN ROUTINE #####
84
85 # Open image and extract pixel data
86 im = Image.open(ARG_IMAGE)
87 xres = im.size[0]
88 yres = im.size[1]
89 # resize to 500px height for convenience
90 im = im.resize((int((float(xres)/yres)*YRES), YRES), Image.BICUBIC)
91 d = list(im.getdata())
92
93 # either the column width or the song length must be fixed width
94 # and if song length is fixed, we have to limit the frequency
95 # spectrum we use to maintain aspect ratio
96 xres = im.size[0]
97 yres = im.size[1]
98 yscale = 22000 / float(yres)
99 # 1/100 of a second of audio for every column in image
100 sampsWithCol = int(SAMPLE_RATE*T_PER_COL)
101
102
103 #because this is easier than finding the flag to disable broadcasting

```

```

104 out = [np.zeros(0), np.zeros(0), np.zeros(0)] #more mehh
105 elfMagic = (float(sampsPerCol)/SAMPLE_RATE)
106 for x in xrange(xres):
107     t = np.linspace(x*elfMagic, (x+1)*elfMagic, num=sampsPerCol)
108     tones = [np.zeros(sampsPerCol), np.zeros(sampsPerCol), np.zeros(sampsPerCol)] # mehh
109     print "{0}: {1}%".format("Color" if ARG_COLOR else "Grayscale",
110                               round(100.0 * x / xres, 2))
111     for y in xrange(yres):
112         p = d[x+xres*y]
113         for c in range(CHANNELS):
114             if p[c] > 10 or p[R] > 10 or p[G] > 10 or p[B] > 10:
115                 if ARG_COLOR:
116                     amplitude = 10**(1-5.25+4.25*(p[c])/(255))
117                 else:
118                     amplitude = 10**(1-5.25+4.25*(p[R]+p[G]+p[B])/(255*3))
119                 tones[c] += oscillator(t, amp=amplitude, freq=yscale * (yres - y))
120     for c in range(CHANNELS):
121         tones[c] = tones[c] + 1
122         tones[c] = tones[c] / math.log(128)
123         out[c] = np.append(out[c], tones[c])
124
125
126 if ARG_COLOR:
127     out = np.array(out)
128     out = out.flatten('F')
129 else:
130     out = out[0]
131
132 writewav(ARG_OUTFILE, CHANNELS, SAMPLE_RATE, 16, int(xres*sampsPerCol), out)

```

6.2 wavToImage.py

```
1  """
2
3  wavToImage.py
4  Creates a spectrogram image from a wav file.
5  Auto-detects color or grayscale.
6
7  Author: Jordan Perr-Sauer, jordan@jperr.com
8          Nicolas Avrutin, nicolasavru@gmail.com
9
10 Revision history:
11     See http://github.com/nicolasavru/DSP-Experiments
12
13 Bugs:
14     - Scaling leaves nasty artifacts
15
16 Todo:
17     - Tweak scaling/encoding algorithm to lessen artifacts
18     - Variable resolution spectrogram
19
20 """
21
22
23 ##### IMPORTS #####
24
25 import wave, sys, struct, math, Image
26 import numpy as np
27
28
29 ##### DEFS AND ARGS #####
30
31 YRES = 400
32 T_PER_COL = 0.03
33
34 ARG_WAVFILE = sys.argv[1]
35 ARG_IMGFILE = sys.argv[2]
36
37
38 ##### MAIN ROUTINE #####
39
40 print "Loading WAV..."
41
42 # Load wav file into array
43 # http://stackoverflow.com/questions/2063284/what-is-the-easiest-way-to-read-wav-files-usi
44 wav = wave.open (ARG_WAVFILE, "r")
45 nchannels, sampwidth, framerate, nframes, comptype, compname = wav.getparams()
46 frames = wav.readframes(nframes * nchannels)
47 out = struct.unpack_from("%dh" % nframes * nchannels, frames)
48 wav.close()
49
50 # Separate loaded frames by channel
51 data = np.zeros((nchannels, nframes), np.int16)
52 for f in xrange(nframes*nchannels):
```



```

53     data[f%nchannels][f/nchannels] = out[f] # integer division used intentionally
54
55 # Compute the dimensions of the encoded image
56 # Setup some constants for decoding
57 sampsPerCol = framerate*T_PER_COL
58 yres = YRES
59 xres = int(math.ceil((nframes)/sampsPerCol))
60
61 # Create a PIL Image
62 if nchannels == 3:
63     im = Image.new("RGB", (xres+1, yres+1))
64 else:
65     im = Image.new("L", (xres+1, yres+1))
66
67
68 print "Generating Spectrogram..."
69
70 # Loop over all "slices" in WAV file.
71 for x in xrange(xres):
72     fft = list()
73     # Perform an FFT on the sound contained in each slice
74     for c in range(nchannels):
75         foo = np.fft.rfft(data[c][x*sampsPerCol:(x+1)*sampsPerCol])
76         fft.append([z.real for z in foo])
77     print "{0}%".format(round(100.0 * x / xres, 2))
78     # Compute pixel colors from fft result
79     for y in xrange(len(fft[0])):
80         if nchannels == 3:
81             c = (int(math.log(abs(fft[0][y])+.001)*10),
82                 int(math.log(abs(fft[1][y])+.001)*10),
83                 int(math.log(abs(fft[2][y])+.001)*10))
84         else:
85             c = int(math.log(abs(fft[0][y])+.001)*10)
86         # Plot pixels, scaling to YRES
87         im.putpixel((int(x), int(yres-int(((float(y)/len(fft[0]))*YRES)))), c)
88
89 print "Encoding Image File..."
90
91 # Save image file using PIL
92 im.save(ARG_IMGFILE, ARG_IMGFILE.split(".")[1].upper())
93
94 print "Done."

```