


ASSIGNMENT 0: DSP-Experiments

FINAL PROJECT

AVRUTIN, PERR-SAUER



CS102

Nicolas Avrutin,

2014 EE

nicolasavru@gmail.com

CS102

Jordan Perr-Sauer

2014 EE

jordan@jperr.com

Wednesday 11th August, 2010

Wednesday 11th August, 2010

Jordan Perr-Sauer

Class of 2014

Electrical Engineering

`jordan@jperr.com`

Nicolas Avrutin

Class of 2014

Electrical Engineering

`nicolasavru@gmail.com`

Stack Overflow

`http://stackoverflow.com/questions/2063284/what-is-the-easiest-way-to-read-wav-files-us`

Scott Wilson, Stanford University

`https://ccrma.stanford.edu/courses/422/projects/WaveFormat/`

1 Description

DSP-Experiments consists of python scripts that can encode images into sound files such that the image can be seen in the spectrogram of the produced sound. The code also provides a method for viewing the spectrogram of the generated audio, decoding the sound back into a (lossy) version of the original image.

In our spectrograms, the x-axis represents time (in the audio file) and the y-axis represents frequency. The brightness of a given pixel represents amplitude (loudness). To encode an image such that it can be viewed in a spectrogram, we first separate the image by it's columns of pixels. Each column of pixels will be represented by a “tone” in the output audio stream

that lasts for a fixed amount of time. These tones consist of sin waves of varying frequencies, at varying amplitudes, depending on the intensity and position of each pixel in the column, respectively. Each column in the input image is converted into such a tone, and these tones are concatenated and output to a wave file.

2 Compilation

You do not need to compile any part of this program. There are some dependencies, however, which you must install prior to execution.

2.1 Dependencies

- Python 2.6 or higher (NumPy and PIL are not yet compatible with Python 3)
- NumPy (<http://numpy.scipy.org/>)
- Python Imaging Library (<http://www.pythonware.com/products/pil/>)

3 Execution

3.1 Manual Page

To produce sound (as a wav file) from an image:

```
python imageToWav.py g|c pathToImage pathToWav
```

The g and c options stand for “Grayscale” or “Color”. Using the c option will result in a 3 channel wav file, where each channel of audio represents one channel of color (red, green, and blue). Using the g option will result in a single channel (mono) wav file and all color data will be lost.

To generate the spectrogram of a produced (or other) wav file:

```
python wavToImage.py pathToWav pathToImage
```

The format of the image will be derived from the extension you give `pathToImage`. You can use any format supported by PIL.

3.2 Sample Inputs

Sample images are provided in the `test_images` dir:

```
python imageToWav.py c test_images/test.png out.wav
```

Sample wav files generate with our script are located in `outputs`:

```
python wavToImage.py outputs/cpu_color.wav out.png
```

4 Features

We support color images by using 3-channel wav files. You can specify grayscale or color using command line arguments at execution. This is discussed in “Execution.”

5 Notes

- Proper functionality is not guaranteed for very large files.
- Scaling and resizing algorithms will be significantly reworked in the future.

6 Listings

6.1 imageToWav.py

```
1  import numpy as np
2  import Image, struct, math, sys
3
4  """
5
6  Usage:
7      ./imageToWav.py [c/b] image_path wav_path
8
9  """
10
11
12  ##### DEFINES AND ARGS #####
13
14  SAMPLE_RATE = 44100
15
16  ARG_IMAGE = sys.argv[2]
17  ARG_OUTFILE = sys.argv[3]
18  ARG_COLOR = False
19  if sys.argv[1] == "c":
20      ARG_COLOR = True
21
22  CHANNELS = 1
23  if ARG_COLOR:
24      CHANNELS = 3
25
26  #rgb aliases
27  R=0
28  G=1
29  B=2
30
31
32  ##### FUNCTIONS #####
33
34  def oscillator(x, freq=1, amp=1, base=0, phase=0):
35      return base + amp * np.sin(2 * np.pi * freq * x + phase)
36
37  def writewav(filename, numChannels, sampleRate, bitsPerSample, time, data):
38      wave = open(filename, 'wb')
39      dataSize = time * sampleRate * numChannels * bitsPerSample / 8
40      #https://ccrma.stanford.edu/courses/422/projects/WaveFormat/
41      ChunkID = 'RIFF'
42      ChunkSize = struct.pack('<I', dataSize + 36)
43      Format = 'WAVE'
44      Subchunk1ID = 'fmt '
45      Subchunk1Size = struct.pack('<I', 16)
46      AudioFormat = struct.pack('<H', 1)
47      NumChannels = struct.pack('<H', numChannels)
48      SampleRate = struct.pack('<I', sampleRate)
49      ByteRate = struct.pack('<I', sampleRate * numChannels * bitsPerSample / 8)
```

```

50 BlockAlign = struct.pack('<H', numChannels * bitsPerSample / 8)
51 BitsPerSample = struct.pack('<H', bitsPerSample)
52 Subchunk2ID = 'data'
53 Subchunk2Size = struct.pack('<I', dataSize)
54 header = ChunkID + ChunkSize + Format + Subchunk1ID + Subchunk1Size + \
55         AudioFormat + NumChannels + SampleRate + ByteRate + BlockAlign + \
56         BitsPerSample + Subchunk2ID + Subchunk2Size
57 wave.write(header)
58 # wav header: 30 s at 44100 Hz, 1 channel of 16 bit signed samples
59 # wave.write('RIFF\x14' + (\x00WAVEfmt \x10\x00\x00\x00\x01\x00\x01\x00D'
60 #           '\xac\x00\x00\x88X\x01\x00\x02\x00\x10\x00data\xfo_(\x00')
61 # little endian
62 # write float64 data as signed int16
63 # amplitude/volume, max value is 32768
64 # higher amplitude causes noise (vertical bars)
65 print "Packing WAV..."
66 (1000 * data).astype(np.int16).tofile(wave)
67 wave.close()
68
69 ##### MAIN ROUTINE #####
70
71 # Open image and extract pixel data
72 im = Image.open(ARG_IMAGE)
73 size = im.size
74 d = list(im.getdata())
75
76
77 xres = size[0]
78 yres = size[1]
79 yscale = 22000 / float(yres)
80 time = int(round(22.0 * xres / yres))
81 xlen = time / float(size[0])
82
83
84 #because this is easier than finding the flag to disable broadcasting
85 out = [np.zeros(0), np.zeros(0), np.zeros(0)] #more mehh
86 for x in xrange(xres):
87     t = np.arange(x*xlen, x*xlen + xlen, 1./SAMPLE_RATE)
88     tones = [np.zeros(t.size), np.zeros(t.size), np.zeros(t.size)] # mehh
89     print "{0}: {1}%".format("Color" if ARG_COLOR else "Grayscale",
90                             round(100.0 * x / xres, 2))
91
92     for y in xrange(yres):
93         p = d[x+xres*y]
94         for c in range(CHANNELS):
95             if p[c] > 10 or p[R] > 10 or p[G] > 10 or p[B] > 10:
96                 if ARG_COLOR:
97                     amplitude = 10**(1-5.25+4.25*(p[c])/(255))
98                 else:
99                     amplitude = 10**(1-5.25+4.25*(p[R]+p[G]+p[B])/(255*3))
100                 tones[c] += oscillator(t, amp=amplitude, freq=yscale * (yres - y))
101     for c in range(CHANNELS):
102         tones[c] = tones[c] + 1
103         tones[c] = tones[c] / math.log(128)

```

```

104         out[c] = np.append(out[c], tones[c])
105
106
107     if ARG_COLOR:
108         out = np.array(out)
109         out = out.flatten('F')
110     else:
111         out = out[0]
112
113     #pad with silence at end if necessary
114     if out.size < SAMPLE_RATE * time * CHANNELS:
115         out = np.append(out, np.zeros(SAMPLE_RATE * time * CHANNELS - out.size))
116
117     writewav(ARG_OUTFILE, CHANNELS, SAMPLE_RATE, 16, time, out)

```

6.2 wavToImage.py

```
1  import wave
2  import sys
3  import struct
4  import math
5
6  import render
7
8  import numpy as np
9
10 xres = 1000
11 yres = 1000
12
13 fname = sys.argv[1]
14
15 # http://stackoverflow.com/questions/2063284/what-is-the-easiest-way-to-read-wav-files-usi
16 wav = wave.open (fname, "r")
17 nchannels, sampwidth, framerate, nframes, comptype, compname = wav.getparams()
18 frames = wav.readframes(nframes * nchannels)
19 out = struct.unpack_from("%dh" % nframes * nchannels, frames)
20 wav.close()
21
22 print nchannels, nframes
23 data = np.zeros((nchannels, nframes), np.int16)
24 for f in xrange(nframes*nchannels):
25     data[f%nchannels][f/nchannels] = out[f] #integer division used intentionally
26
27 # set this to the number of the channel you want to use
28 channel = 0
29 #loadWav(fname)
30 time = float(nframes) / framerate
31 print time, framerate
32 yres = int(22 * xres / time)
33 print yres
34 #yscale = float(yres) / 22000
35
36
37 interval = nframes / xres
38 print interval, nframes
39
40 screen = render.createScreen(xres, yres)
41
42 def generateColor(val):
43     v = math.log(abs(val)+.001)*10
44
45 for x in range(xres):
46     fft0 = np.fft.rfft(data[0][x*interval:(x+1)*interval+10])
47     fft1 = np.fft.rfft(data[1][x*interval:(x+1)*interval+10])
48     fft2 = np.fft.rfft(data[2][x*interval:(x+1)*interval+10])
49     fft0 = [z.real for z in fft0]
50     fft1 = [z.real for z in fft1]
51     fft2 = [z.real for z in fft2]
52     # print len(fft)
```



```

53 #     print fft
54     print "{0}%".format(round(100.0 * x / xres, 2))
55     for y in range(interval / 2):
56         c = [math.log(abs(fft0[y])+.001)*10,
57              math.log(abs(fft1[y])+.001)*10,
58              math.log(abs(fft2[y])+.001)*10]
59     #     print x, y, c
60     render.plot(x, yres-y, c, screen)
61
62 #render.display(screen)
63 render.saveExtension(screen, sys.argv[2])

```